




#1

Q1 · HOW DID YOU APPROACH THE CHALLENGE?

A LOOP THAT CHANGES ITS MIND



ONE REAL LAP · THE DUELIST

FIRST IDEA: LET THE LLM NEGOTIATE

- 1** Sat · Duels II With the LLM deciding, decisions took **9.4 s**, 29 % over 10 s. Sunday's tick: **15 s**. **V**
- 2** Sat 22:30 **Code decides** price, day and accept; the LLM only writes. Checked first on a simulator validated on Duels II.
- 3** Sat 22:48 → Sun 09:45 Branch **duelist-loop**: Haiku ≈ 2 s, guards, 8 s failover, hot-reloaded params.
- 4** Sun 10:59 Duels III, shorter clock, harder decay: **57 of 68** (84 %); the Grand Final **27 of 34** (79 %), both above the field. The monitor reviews every deal. **V**
- 5** 11:47 → 12:37 Review + Duel Lab find a missed last-tick accept (11124); fix shipped (f57a002: last-tick accept + worth floor), restarted on the **Duel Lab's GO**.

KEPT THE LLM'S WORDS. REPLACED ITS DECISIONS.

HOW WE STAYED
IN SYNC

Git as the bus
one file per owner

directives.md
big calls + guardrails

Session messages
Claude ↔ Claude

STATUS.md
auto, every 5 min

Phone alerts
to Lucas + Dani

Dani's dashboard
live, read-only

Neon hub
one copy of the game

Q2 · WHAT DID YOU BUILD? · Q3 · WHY THAT WAY?

WHO DECIDES, WHO ACTS, WHO CHECKS

→ AUTHORITY / GAME WRITE - - - → DATA

DECIDE

- LUCAS**
strategy · goals + hard limits
- ALEKS**
wrote the duelist + the hub
- DANI**
the room · rivals · dashboard

CHIEF OF STAFF
Claude Code · makes the big calls, logs them with evidence · **never writes to the game**

DIRECTIVES + LIMITS

ACT · TWO TRADING ROLES

OPERATOR
Claude Code · the only trade writer

trader · taker book · maker
dealer bots closers

DUELIST
code decides · Haiku writes the words

guards 8 s failover
hot-reloaded params
C → A switch

Sunday duelist: branch duelist-loop @ f57a002

WRITES

WRITES

SENSE · CHECK

MARKET
venue v10 at 0 % · Market Test recorder

COLLECTOR + STATUS
feed · me · board → data/, STATUS.md


DASHBOARD + HUB
Dani's screen · Neon archive

ANALYSTS
scout + judge on the Claude API

DUEL LAB
simulator · tests params first

VERIFIERS
fresh subagents · contrarian reviews

READS



● Live · 14:26
DANI'S DASHBOARD
the room's view: us #1, every team's race, read-only

THE BAZAAR · GAME API

WRITE offers · accepts · threads · duels READ me · clock · board · event stream

ONE TEAM KEY · 5 REQ/S
1 ACCEPT PER TICK

WHO DECIDES, WHO ACTS, WHO CHECKS

→ AUTHORITY / GAME WRITE - - - → DATA

DECIDE

1

CODE DECIDES, THE LLM WRITES

The duelist's price, day and accept are code; Haiku only writes the words, in ≈ 2 s instead of 9.4 s.

Guards in code, not in the prompt.

BECAUSE A 15 S TICK DOESN'T WAIT

2

ONE TRADER PER JOB, LIMITS IN WRITING

Only the Operator and the Duelist trade; the Market session only posts our venue's ads.

Lucas changes a hard limit with a **GUARDRAIL** line in directives.md.

BECAUSE ONE KEY, 5 REQ/S AND 1 ACCEPT PER TICK ARE SHARED

3

CHECK BEFORE YOU ACT

A fresh verifier and contrarian reviews on the big calls; the Duel Lab simulates before every Sunday duelist start.

BECAUSE THE AUTHOR CAN'T GRADE ITS OWN WORK

Sunday duelist: branch duelist-loop @ f57a002

WRITES

WRITES

READS

THE BAZAAR · GAME API

WRITE offers · accepts · threads · duels READ me · clock · board · event stream

ONE TEAM KEY · 5 REQ/S
1 ACCEPT PER TICK

WHAT WE **LEARNED**

TECHNICAL DRAFT

Code decides, the LLM writes. On a 15 s tick, latency is a rule: Haiku $\approx$ 2 s vs 9.4 s when the LLM decided.

The generator can't grade itself. A separate verifier caught errors the author missed, in this very deck too.

NEGOTIATION DRAFT

Decay is per exchange, so silence is free. Hold instead of talking.

Take the good offer when it's there. 23 Duels III deals closed below an offer the rival had already made ($\approx$ 197 P, the monitor's price check).

MARKETPLACE DRAFT

Value created, not activity. A trade with a rival lifts the rival as much as you: CHA-11, +1.48 each.

Pushing partners moved trades; the bounty didn't. 0 fills at 10:45, 8 at the close; the bounty paid nobody.

Nobody beat the free stall. 0.5 in all eleven Market Tests: no bond needed.

WITH ONE MORE DAY DRAFT

Merge the Sunday duelist into main

Switch on the broker we staged

A verifier gate on every directive

THREE HUMANS SET THE LIMITS. CLAUDE RAN THE BAZAAR, AND CHECKED ITSELF.

TEAM 5

3 HUMANS


7+ CLAUDE CODE SESSIONS

19 DAEMONS BUILT

2 TRADING ROLES

FROM #6 TO #1

Friday final #6 · first #1 Sat 16:44 · #1 since 12:35 today · board 37.73, lead 1.97 at 14:56 v



FIVE BELIEFS OUR DATA BROKE

	WE BELIEVED	WHAT BROKE IT	WHAT WE CHANGED
Dealers	Any trade at a gain scores, dealers included.	Dealer gains clip to 0, losses count in full: RET-09 at 87, worth 77 → -10 . v	Pages close only through team trades. Dealers only for the ladder; autoflip killed.
Duels	An LLM should negotiate the duels.	LLM decisions took 9.4 s on average, 29 % over 10 s ; Sunday's tick: 15 s. v	Code decides, Haiku only writes. 8 s failover. Duels III: 57 of 68 closed. v
Scoring	Trade gains cap at 50 per round (10:31).	A 5 P sale moved neg_points 50.0 → 52.2 at 12:17. v	Corrected at 12:22. Buyers back on; the next +50 trade, CHA-11, added +1.48 . v
Rivals	A trade is good if we gain.	CHA-11 from Team 10: we gained +1.48, so did Team 10 . v	Lock the lead (13:12): zero trades with the five closest rivals.
Market	A 0 % fee venue attracts trades.	v10 at 0 %: 0 fills today by 10:45. v	We pushed pairs and partners onto v10: our RET-01 pair, Team 13's bot, Team 10's listings. 8 fills by the close . v

WHERE THE POINTS CAME FROM

37.26

#1 · LEAD 3.09 · 13:20 V

Game points at 13:20 (tick 2422) by the part of the system that earned them; final board 37.73 · V measured L modelled



Scored 0: RET-11 at 240 · MAL page · cash held v10: the seller bounty paid nobody Not shown: judges' 40 · Grand Final

WHAT WE BUILT **AROUND** THE BAZAAR

A MARKET WITH MORE THAN TWO SIDES



THE CLUB MATCHMAKER

We estimate **every team's album from the public feed** and shared want-lists, match one team's spare copy to another team's missing card, page-closers first, and price each deal **between the two values**, so both sides gain.

Six teams invited to a club, with hosting **designed to rotate across members' markets**, so the matching pays more than one venue.

`tools/matchmaker.py` · `intel/club-pitch.md` · `brokered RET-01 Team 2` → `Team 8 on v10`



DEMAND MODEL + HUB

What **every card is worth to every team**, from public behaviour alone, no LLM. A shared Neon archive keeps the history the 500-event feed drops.

`hub/demand.py` · `hub/README.md`



MARKET TEST LAB

A recorder captures **every bench session**; an offline simulator tests broker strategies against the free stall before we risk a bond.

`broker/record_bench.py` · `broker/sim.py`



LIVE REACTOR

Reads the public **event stream** without our key, flags every ask we can take at a gain, and **pages the Operator** and our phones.

`tools/reactor.py` · `tools/notify.py`



JUDGES' SHOWCASE

Dani's live dashboard and a **/show** page: read-only, built to run offline, the whole game on one screen.

`dashboard/server.py` · `dashboard/show.html`